

**CSC 370 - SPRING 2018
DATABASE SYSTEMS
ASSIGNMENT 2
UNIVERSITY OF VICTORIA**

Due: Wednesday, Feb. 28th, 2018 at 11:55pm. **Late assignments will not be accepted.**

This assignment will be accepted electronically. See the ‘Submission and Evaluation’ section below for details on the submission process and expected formatting of your answers. For all of the questions below, your answer must be **one** SQL query (including a terminating semicolon) which runs without errors on the `studdb1.csc.uvic.ca` or `studdb2.csc.uvic.ca` PostgreSQL database servers. Note that timeout errors (in which the server terminates your query for exceeding the maximum execution time) are considered errors. Queries which have errors will receive a mark of zero. All queries without errors will be marked out of two, with full marks given only to queries which produce the correct output and contain no assumptions besides the data given in the question (see the advice sections below for more details).

Question 1: IMDB Queries [18 marks]

Create queries for each of the data retrieval problems below, using the `imdb` database. Place your answers in the appropriate positions in the `a2q1_queries.txt` file. In the questions below, any reference to ‘films’ refers to titles with `title_type = 'movie'`.

Hint: When you need the ‘primary name’ of a title, add the following to the `WITH` clause of your query and then join the resulting `primary_names` subquery to `titles` as needed.

```
primary_names as (select title_id, name as primary_name
                    from title_names where is_primary = true)
```

For example, the query below uses the ‘primary name’ to find the production year, title ID and duration of all films (`title_type = 'movie'`) with primary title ‘The Shining’.

with

```
primary_names as (select title_id, name as primary_name
                    from title_names where is_primary = true)
```

```
select *
```

```
from
```

```
    titles
```

```
  natural join
```

```
    primary_names
```

```
where primary_name = 'The Shining' and title_type = 'movie';
```

The result of the above query is shown below.

Query Result				
title_id	title_type	year	length_minutes	primary_name
6812278	movie	2017	136	The Shining
81505	movie	1980	146	The Shining

- (a) Find the primary name, year and title ID of all titles from the year 1989 with a length of 180 minutes and a title type of 'tvSpecial'.

Expected Query Result		
primary_name	year	title_id
Survivor Series	1989	264059
Starrcade	1989	348114
The 16th Annual American Music Awards	1989	790592
The 1989 Miss Tennessee Pageant	1989	1837666

- (b) Find the primary name, year and length (in minutes) of all films whose total length in minutes is at least 4320 (72 hours).

Expected Query Result		
primary_name	year	length_minutes
Modern Times Forever	2011	14400
Beijing 2003	2004	9000
Nari	2017	6017
Hunger!	2015	6000
London EC1	2015	5460
The Cure for Insomnia	1987	5220
Ember Glow	2015	4980
Fail	2016	4680
Writing on Snow	2017	4320

- (c) Find the primary name, year and length (in minutes) of all films which have 'Meryl Streep' in the cast/crew **and** have a production year of 1985 or earlier. Note: Use the `cast_crew` table, not the `known_for` table.

Expected Query Result		
primary_name	year	length_minutes
Kramer vs. Kramer	1979	105
The Seduction of Joe Tynan	1979	108
The French Lieutenant's Woman	1981	124
Sophie's Choice	1982	150
Still of the Night	1982	93
Silkwood	1983	131
Falling in Love	1984	106
Out of Africa	1985	161
Plenty	1985	121

- (d) Find the primary name, year and length (in minutes) of all films which are associated with **both** of the genres 'Film-Noir' and 'Action'. Use the `title_genres` table to map titles to their genres (titles may have any number of genres).

Expected Query Result		
primary_name	year	length_minutes
Blackmail	1947	67
Dangerous Mission	1954	75
Dick Tracy	1945	61
Dick Tracy vs. Cueball	1946	62
His Kind of Woman	1951	120
Peking Express	1951	95
Road House	1948	95
Rogues' Regiment	1948	86
Scotland Yard Investigator	1945	68
Sirocco	1951	98
The Pay Off	1942	74
Unmasked	1950	60

- (e) Find the names of all people who were associated as cast or crew (using the `cast_crew` table) with the film 'The Big Lebowski'. Note that there is exactly one film (with `title_type = 'movie'`) in the database with the name 'The Big Lebowski' (but there may be other titles, like TV episodes, with that name).

Expected Query Result	
name	
Carter Burwell	
Ethan Coen	
Jeff Bridges	
Joel Coen	
John Goodman	
Julianne Moore	
Rick Heinrichs	
Roger Deakins	
Steve Buscemi	
Tricia Cooke	

- (f) Find the names of all people who were associated as writers or directors (using the appropriate relationship tables) with the movie 'Die Hard'. Again, there is only one title in the database which is both a movie and has primary name 'Die Hard'.

Expected Query Result	
name	
Jeb Stuart	
John McTiernan	
Roderick Thorp	
Steven E. de Souza	

- (g) Find the primary name and length (in minutes) of all films that 'Tom Cruise' is known for (using the `known_for` table, not the `cast_crew` table).

Expected Query Result	
primary_name	length_minutes
Jerry Maguire	139
Minority Report	145
The Last Samurai	154
Top Gun	110

- (h) Find the primary name, year and length (in minutes) of all films which have **both** ‘Tom Hanks’ and ‘Meryl Streep’ as cast/crew. Use the `cast_crew` table, not the `known_for` table.

Expected Query Result		
primary_name	year	length_minutes
Everything Is Copy	2015	89
The Ant Bully	2006	88
The Post	2017	115

- (i) Find the primary name and year of all films which were directed by ‘Steven Spielberg’ and are associated with the genre ‘Thriller’. Use the `title_genres` table to map titles to their genres.

Expected Query Result	
primary_name	year
Bridge of Spies	2015
Firelight	1964
Jaws	1975
Jurassic Park	1993
Munich	2005
War of the Worlds	2005

Question 2: BC Ferries Queries [10 marks]

The queries you write below should work correctly on any of the BC Ferries databases (`ferries_1month`, `ferries_3months`, `ferries_6months`, `ferries_9months` or `ferries_12months`). For comparison, sample output is shown for both `ferries_1month` and `ferries_12months`.

- (a) Print the names of all vessels which have served on route number 1. Vessel names must not appear more than once.

Expected Query Result (<code>ferries_1month</code>)	
vessel_name	
Coastal Celebration	
Coastal Renaissance	
Queen of New Westminster	
Spirit of Vancouver Island	

Expected Query Result (ferries_12months)
vessel_name
Coastal Celebration
Coastal Inspiration
Coastal Renaissance
Queen of New Westminster
Spirit of British Columbia
Spirit of Vancouver Island

- (b) Print the total number of sailings per vessel in the database, showing only those vessels with at least one sailing.

Expected Query Result (ferries_1month)	
vessel_name	count
Bowen Queen	334
Coastal Celebration	242
Coastal Inspiration	209
Coastal Renaissance	216
Queen of Alberni	223
Queen of Capilano	118
Queen of Coquitlam	177
Queen of Cowichan	286
Queen of New Westminster	64
Queen of Oak Bay	89
Queen of Surrey	430
Skeena Queen	240
Spirit of Vancouver Island	88

Expected Query Result (ferries_12months)	
vessel_name	count
Bowen Queen	411
Coastal Celebration	2342
Coastal Inspiration	1983
Coastal Renaissance	2173
Island Sky	205
Mayne Queen	1
Queen of Alberni	2255
Queen of Capilano	4780
Queen of Coquitlam	2277
Queen of Cowichan	2339
Queen of Cumberland	18
Queen of New Westminster	1382
Queen of Oak Bay	2571
Queen of Surrey	4853
Skeena Queen	2724
Spirit of British Columbia	1425
Spirit of Vancouver Island	2203

- (c) Print the names and number of routes served for all vessels which served at least 2 routes.

Expected Query Result (ferries_1month)	
vessel_name	num_routes
Bowen Queen	2
Coastal Renaissance	2
Queen of Coquitlam	2
Queen of Cowichan	2
Queen of Surrey	2

Expected Query Result (ferries_12months)	
vessel_name	num_routes
Bowen Queen	3
Coastal Inspiration	2
Coastal Renaissance	3
Queen of Coquitlam	3
Queen of Cowichan	3
Queen of New Westminster	2
Queen of Surrey	2

- (d) For each route which appears at least once in the sailings table, print the route number and the name and production year of the oldest vessel(s) on the route. Note that there may be multiple vessels tied for oldest (each should be printed separately). Hint: You will likely need a join on a subquery.

Expected Query Result (ferries_1month)		
route_number	vessel_name	year_built
1	Queen of New Westminster	1964
2	Queen of Coquitlam	1976
2	Queen of Cowichan	1976
3	Bowen Queen	1965
4	Skeena Queen	1997
8	Bowen Queen	1965
30	Queen of Alberni	1976

Expected Query Result (ferries_12months)		
route_number	vessel_name	year_built
1	Queen of New Westminster	1964
2	Queen of Coquitlam	1976
2	Queen of Cowichan	1976
3	Bowen Queen	1965
4	Bowen Queen	1965
4	Mayne Queen	1965
8	Bowen Queen	1965
30	Queen of New Westminster	1964

- (e) List all vessels which used any port (source or destination) that was at any time used (as either source or destination) by the vessel 'Queen of New Westminster'. The result should contain the Queen of New Westminster itself. Remember not to make any assumptions about the data.

Expected Query Result (ferries_1month)	
vessel_name	
Coastal Celebration	
Coastal Inspiration	
Coastal Renaissance	
Queen of Alberni	
Queen of New Westminster	
Skeena Queen	
Spirit of Vancouver Island	

Expected Query Result (ferries_12months)
vessel_name
Bowen Queen
Coastal Celebration
Coastal Inspiration
Coastal Renaissance
Mayne Queen
Queen of Alberni
Queen of Coquitlam
Queen of Cowichan
Queen of Cumberland
Queen of New Westminster
Skeena Queen
Spirit of British Columbia
Spirit of Vancouver Island

Advice: Assume Nothing

You will lose marks if your query contains any ‘hard-coded’ assumptions about the data in the database other than the data given in the question. For example, suppose you were asked to create the following query.

Using the **fruit** database, print the first and last names of every customer who placed an order containing the product ‘Pear’, along with the order number of their order.

One query which correctly implements the requirements above (and would receive full marks) is

```
select distinct customer_firstname, customer_lastname, order_num
from
  products
natural join
  orders
natural join
  order_contents
where products.name = 'Pear'
order by order_num asc;
```

The result of the above query is shown below.

Query Result		
customer_firstname	customer_lastname	order_num
Franz	Kafka	1001
Fiona	Framboise	1002

It is possible to write plenty of other queries which produce the same result (and any such query which makes no assumptions about the data would receive full marks). In some cases, it is possible to write a query which takes advantage of properties of the data which cannot generally be assumed. For example, if you happen to know that the product ID number for ‘Pear’ is 2, the query below would produce the same result as above.

Bad Query 1:

```
select distinct customer_firstname, customer_lastname, order_num
  from
    orders
  natural join
    order_contents
 where product_id = 2
 order by order_num asc;
```

Taking the assumptions even further, if you happen to know that the only orders meeting the criteria of the question are orders 1001 and 1002, the query below would produce the correct output as well.

Bad Query 2:

```
select distinct customer_firstname, customer_lastname, order_num
  from
    orders
 where order_num = 1001 or order_num = 1002
 order by order_num asc;
```

Both of the queries above would lose marks (and the second bad query would likely receive no marks at all), since both include hard-coded assumptions about the data in the database. If the product ID of **Pear** were to change, or more orders were added to the database, the two bad queries would no longer work (but the original, correct query would continue to work properly).

Advice: Don't Plagiarize

You are encouraged to discuss solution methods with your peers, and even to look up possible solution ideas on the internet, but all of your submitted queries must be your own work. As a rule of thumb, to ensure you do not accidentally plagiarize, do not look at anyone else's queries (or allow them to see yours). Additionally, if you use a version control system (such as Github) to store your work, ensure that the repository is private. If your queries are posted in a public setting (even inadvertently), you may become entangled in any ensuing academic integrity investigation if your work is copied by someone else.

Submission and Evaluation

This assignment will be marked through a combination of automated testing (that is, running your submitted queries and examining the result) and human inspection. To expedite the marking process (and ensure consistency between different student submissions), you are required to submit your queries for each question inside of premade query template files. Two empty template files have been posted to **conneX: a2q1_queries.txt** and **a2q2_queries.txt**. Please place your query for each subquestion in the indicated spaces in the templates before submitting (we will be using an automated system to extract each query individually for marking, so failure to comply with this requirement will likely result in you receiving a mark of zero for any queries which do not meet the

formatting requirements). Notice that the provided files have the `.txt` extension instead of `.sql` (which would normally be used for SQL queries); this is due to a technical limitation of `conneX` (which does not properly handle files with the `.sql` extension when they are submitted, probably because it sees them as a security risk).

You are required to submit your answers to each question in two files called `a2q1_queries.txt` and `a2q2_queries.txt`, using the provided templates as a starting point. Your answer for each query must consist of a single SQL statement (which may have a `WITH` clause containing multiple subqueries, and/or several `SELECT` statements joined by set operators). As a point of reference, your answer is a ‘single SQL statement’ if it contains only one semicolon (at the end of the query). Although your files will be associated with your account, please ensure that each submitted file contains a comment with your name and student number.

You are permitted to delete and resubmit your submission as many times as you want before the due date, but no submissions or resubmissions will be accepted after the due date has passed. You will receive a mark of zero if you have not officially submitted your assignment (and received a confirmation email) before the due date.

Only the files that you submit through `conneX` will be marked. The best way to make sure your submission is correct is to download it from `conneX` after submitting and test it. You are not permitted to revise your submission after the due date, and late submissions will not be accepted, so you should ensure that you have submitted the correct version of your code before the due date. `conneX` will allow you to change your submission before the due date if you notice a mistake. After submitting your assignment, `conneX` will automatically send you a confirmation email. **If you do not receive such an email, you did not submit the assignment.** If you have problems with the submission process, send an email to the instructor **before** the due date.